

Assignment1

September 17, 2023

1 Assignment 1

1.1 Author Details

Name: Yelisetty Karthikeya S M

Roll No: 21CS30060

1.2 Importing Libraries

```
[1]: # Base libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tabulate import tabulate
from tqdm import tqdm
from pickle import dump

# Preprocessing
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE

# Models
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier

# Evaluation
from sklearn.model_selection import cross_validate
from sklearn.metrics import classification_report, accuracy_score, \
    precision_score, recall_score, f1_score, roc_auc_score
```

```
# Initializations
%matplotlib inline
sns.set_style('darkgrid')
pd.set_option('display.max_columns', None)
plt.style.use('ggplot')
```

1.3 Data Preprocessing and Exploratory Data Analysis

1.3.1 Loading dataset

```
[2]: df = pd.read_csv("./heart_disease_prediction/framingham.csv")
```

1.3.2 Sample data

```
[3]: df.head()
```

```
[3]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	\
0	1	39	4.0	0	0.0	0.0	0	
1	0	46	2.0	0	0.0	0.0	0	
2	1	48	1.0	1	20.0	0.0	0	
3	0	61	3.0	1	30.0	0.0	0	
4	0	46	3.0	1	23.0	0.0	0	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0	0	0	195.0	106.0	70.0	26.97	80.0	77.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	76.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	70.0	
3	1	0	225.0	150.0	95.0	28.58	65.0	103.0	
4	0	0	285.0	130.0	84.0	23.10	85.0	85.0	

	TenYearCHD
0	0
1	0
2	0
3	1
4	0

1.3.3 Information about the dataset

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
#   ...
```

```

---  -----
0  male          4240 non-null  int64
1  age           4240 non-null  int64
2  education     4135 non-null  float64
3  currentSmoker 4240 non-null  int64
4  cigsPerDay    4211 non-null  float64
5  BPMeds        4187 non-null  float64
6  prevalentStroke 4240 non-null  int64
7  prevalentHyp  4240 non-null  int64
8  diabetes      4240 non-null  int64
9  totChol       4190 non-null  float64
10 sysBP         4240 non-null  float64
11 diaBP         4240 non-null  float64
12 BMI           4221 non-null  float64
13 heartRate     4239 non-null  float64
14 glucose       3852 non-null  float64
15 TenYearCHD    4240 non-null  int64

```

dtypes: float64(9), int64(7)

memory usage: 530.1 KB

```
[5]: df.describe()
```

```

[5]:
count      male          age  education  currentSmoker  cigsPerDay \
mean      0.429245    49.580189    1.979444    0.494104    9.005937
std        0.495027     8.572942    1.019791    0.500024   11.922462
min         0.000000    32.000000    1.000000    0.000000    0.000000
25%         0.000000    42.000000    1.000000    0.000000    0.000000
50%         0.000000    49.000000    2.000000    0.000000    0.000000
75%         1.000000    56.000000    3.000000    1.000000   20.000000
max         1.000000    70.000000    4.000000    1.000000   70.000000

count      BPMeds  prevalentStroke  prevalentHyp    diabetes    totChol \
mean      0.029615     0.005896     0.310613     0.025708   236.699523
std        0.169544     0.076569     0.462799     0.158280   44.591284
min         0.000000     0.000000     0.000000     0.000000   107.000000
25%         0.000000     0.000000     0.000000     0.000000   206.000000
50%         0.000000     0.000000     0.000000     0.000000   234.000000
75%         0.000000     0.000000     1.000000     0.000000   263.000000
max         1.000000     1.000000     1.000000     1.000000   696.000000

count      sysBP      diaBP      BMI    heartRate    glucose \
mean     132.354599    82.897759    25.800801    75.878981    81.963655
std       22.033300    11.910394     4.079840    12.025348    23.954335
min       83.500000    48.000000    15.540000    44.000000    40.000000

```

25%	117.000000	75.000000	23.070000	68.000000	71.000000
50%	128.000000	82.000000	25.400000	75.000000	78.000000
75%	144.000000	90.000000	28.040000	83.000000	87.000000
max	295.000000	142.500000	56.800000	143.000000	394.000000

	TenYearCHD
count	4240.000000
mean	0.151887
std	0.358953
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

1.3.4 Null Values inspection

```
[6]: df.isnull().sum()
```

```
[6]: male                0
     age                0
     education          105
     currentSmoker       0
     cigsPerDay          29
     BPMeds              53
     prevalentStroke     0
     prevalentHyp        0
     diabetes            0
     totChol             50
     sysBP               0
     diaBP               0
     BMI                 19
     heartRate           1
     glucose             388
     TenYearCHD          0
     dtype: int64
```

Simple solution is to drop the rows with missing values. But this is not a good solution as we might lose some important information. So, we will use imputation to fill in the missing values using the mode of the column for nominal values and mean for continuous values.

```
[7]: for col in df.columns:
      df[col].fillna(df[col].mode()[0], inplace=True)
```

```
[8]: df.isnull().sum()
```

```
[8]: male          0
      age          0
      education    0
      currentSmoker 0
      cigsPerDay    0
      BPMeds       0
      prevalentStroke 0
      prevalentHyp  0
      diabetes     0
      totChol      0
      sysBP        0
      diaBP        0
      BMI          0
      heartRate    0
      glucose      0
      TenYearCHD   0
      dtype: int64
```

1.3.5 Changing the data types

```
[9]: df.dtypes
```

```
[9]: male          int64
      age          int64
      education    float64
      currentSmoker int64
      cigsPerDay    float64
      BPMeds       float64
      prevalentStroke int64
      prevalentHyp  int64
      diabetes     int64
      totChol      float64
      sysBP        float64
      diaBP        float64
      BMI          float64
      heartRate    float64
      glucose      float64
      TenYearCHD   int64
      dtype: object
```

These are the specifications according to the dataset description:

```
[10]: nominals = ['male', 'currentSmoker', 'BPMeds', 'prevalentStroke',
                  ↪ 'prevalentHyp', 'diabetes', 'TenYearCHD']
      continuous = ['age', 'education', 'cigsPerDay', 'totChol', 'sysBP', 'diaBP',
                  ↪ 'BMI', 'heartRate', 'glucose']
```

```
df[nominals] = df[nominals].astype('int64')
df[continuous] = df[continuous].astype('float64')
```

```
[11]: df.dtypes
```

```
[11]: male                int64
age                  float64
education            float64
currentSmoker        int64
cigsPerDay           float64
BPMeds               int64
prevalentStroke       int64
prevalentHyp          int64
diabetes              int64
totChol              float64
sysBP                float64
diaBP                float64
BMI                  float64
heartRate            float64
glucose              float64
TenYearCHD           int64
dtype: object
```

1.3.6 Dividing the dataset into features and target

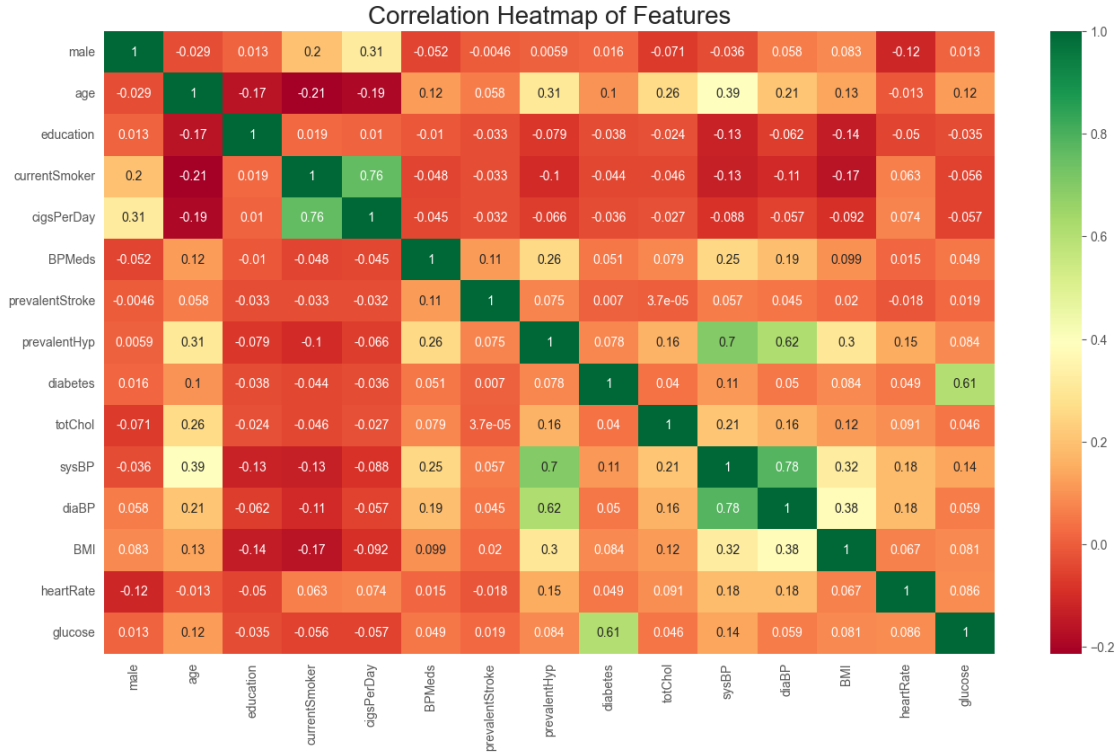
```
[12]: df_features = df.drop('TenYearCHD', axis=1)
df_target = df['TenYearCHD']
```

1.3.7 Normalizing the data

```
[13]: scale = StandardScaler()
df_features[continuous] = scale.fit_transform(df_features[continuous])
```

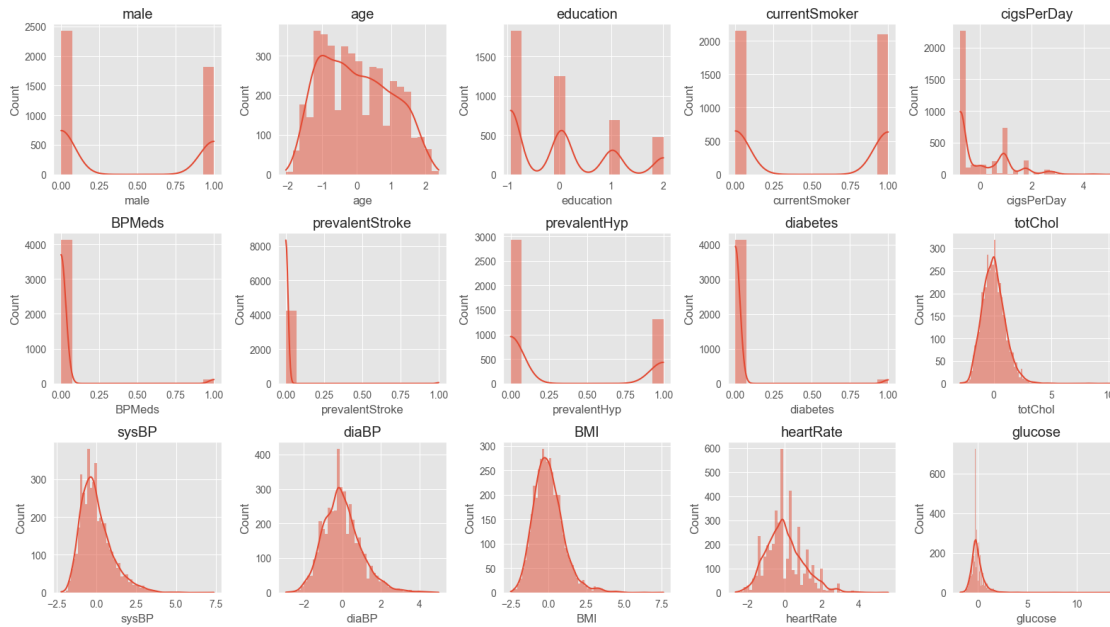
1.3.8 Correlation Heatmap

```
[14]: plt.figure(figsize=(16, 9))
sns.heatmap(df_features.corr(), annot=True, cmap="RdYlGn")
plt.title("Correlation Heatmap of Features", fontsize=20)
plt.show()
```



1.3.9 Distribution of individual features

```
[15]: # Distribution of individual features
plt.figure(figsize=(16, 9))
for i, col in enumerate(df_features.columns):
    plt.subplot(3, 5, i+1)
    sns.histplot(df_features[col], kde=True, linewidth=0)
    plt.title(col)
plt.tight_layout()
plt.show()
```



1.3.10 Class Distribution

```
[16]: values = df_target.value_counts()
indices = [k for k in values.index]
plt.figure(figsize=(10,4))
plt.title('Class distribution in test set')
plt.barh(indices, values, color=['green', 'red'], label=['No disease', 'Disease'], height=0.5)
plt.xlabel('Frequency')
plt.ylabel('Class')
plt.yticks(indices, ['No disease', 'Disease'])
plt.legend()
plt.show()
```




1.3.11 Function to load and preprocess the data


```
[17]: def load_data():
    df = pd.read_csv("./heart_disease_prediction/framingham.csv")
    nominals = ['male', 'currentSmoker', 'BPMeds', 'prevalentStroke',
    ↪ 'prevalentHyp', 'diabetes', 'TenYearCHD']
    continuous = ['age', 'education', 'cigsPerDay', 'totChol', 'sysBP',
    ↪ 'diaBP', 'BMI', 'heartRate', 'glucose']
    for col in df.columns:
        df[col].fillna(df[col].mode()[0], inplace=True)
    df[nominals] = df[nominals].astype('int64')
    df[continuous] = df[continuous].astype('float64')
    df_features = df.drop('TenYearCHD', axis=1)
    df_target = df['TenYearCHD']
    scale = StandardScaler()
    df_features[continuous] = scale.fit_transform(df_features[continuous])
    return df_features, df_target
```

1.4 Feature Selection and Scaling

```
[18]: # PCA
pca = PCA()
pca.fit(df_features)
df_pca = pca.transform(df_features)
```

1.4.1 Number of Components vs. Variance Explained

```
[19]: plt.figure(figsize=(15, 5))
plt.plot(np.cumsum(pca.explained_variance_ratio_)*100)
plt.xlabel('Number of Components')
plt.ylabel('Variance (%)')
plt.title('Explained Variance')
plt.grid(visible=True)
plt.show()
```



1.4.2 Dataset after PCA

As the elbow point is at 7, we will reduce the number of features to 7.

```
[20]: number_of_components = 7
df_features = pd.DataFrame(df_pca[:, :number_of_components], columns=[f'PC{i}']
    ↪ for i in range(1, number_of_components+1)])
df_features.head()
```

```
[20]:
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
0	-1.962665	-0.306698	-1.240161	1.541623	0.783884	-0.377466	0.881480
1	0.125056	-0.014960	0.090382	0.553265	0.502263	-1.724939	0.985375
2	-0.488438	0.799507	0.491221	-1.150302	-0.407634	0.038308	0.087517
3	1.273526	0.929873	-0.106130	0.309940	-0.382192	2.370008	-0.456905
4	-0.397436	1.362755	0.597880	1.248440	-1.026758	-0.033045	0.159698

```
[21]: df_features.info()
```

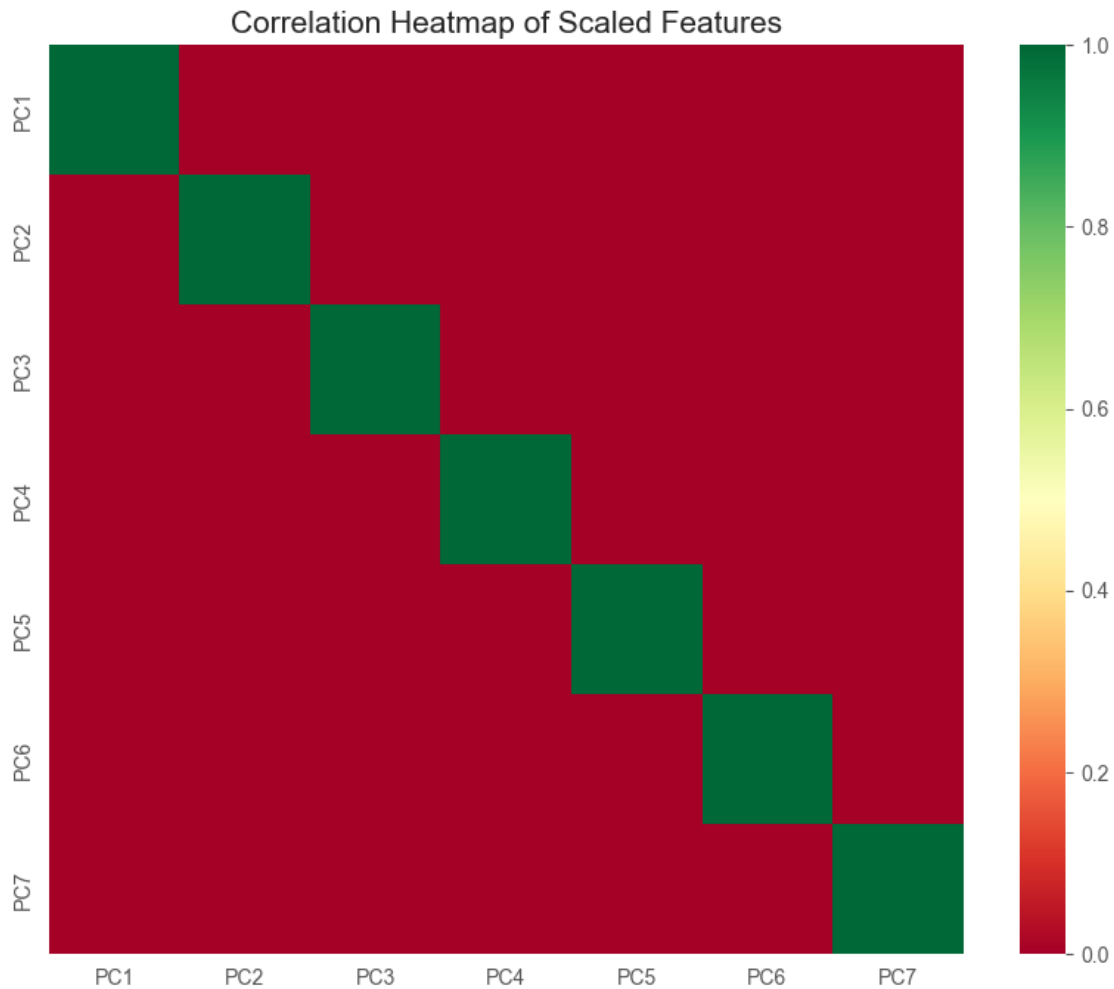
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 7 columns):
#   Column  Non-Null Count  Dtype
---  -

```

```
0  PC1      4240 non-null  float64
1  PC2      4240 non-null  float64
2  PC3      4240 non-null  float64
3  PC4      4240 non-null  float64
4  PC5      4240 non-null  float64
5  PC6      4240 non-null  float64
6  PC7      4240 non-null  float64
dtypes: float64(7)
memory usage: 232.0 KB
```

1.4.3 Correlation Heatmap after PCA

```
[22]: plt.figure(figsize=(10, 8))
sns.heatmap(df_features.corr(), annot=False, cmap="RdYlGn")
plt.title("Correlation Heatmap of Scaled Features", fontsize=15)
plt.show()
```



1.4.4 Function to perform feature selection and scaling

```
[23]: # We selected number of components as 7, because it explains 92% of the
      ↪variance and is an elbow point in the graph.
def feature_selection(df, number_of_components: int = 7):
    # df is df_features in our case
    pca = PCA()
    pca.fit(df)
    df_pca = pca.transform(df)
    df = pd.DataFrame(df_pca[:, :number_of_components], columns=[f'PC{i}' for i
    ↪in range(1, number_of_components+1)])
    return df
```

1.5 Splitting the data into train and test sets

```
[24]: X_train, X_test, y_train, y_test = train_test_split(df_features, df_target,
      ↪test_size=0.2, random_state=42)
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (3392, 7)
X_test shape: (848, 7)
y_train shape: (3392,)
y_test shape: (848,)
```

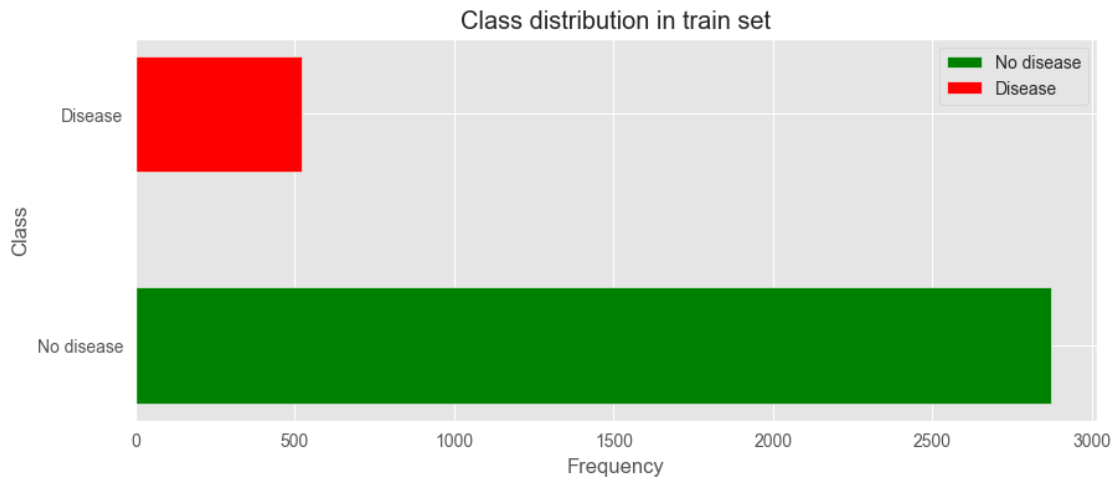
```
[25]: df_target.value_counts()
```

```
[25]: TenYearCHD
0      3596
1       644
Name: count, dtype: int64
```

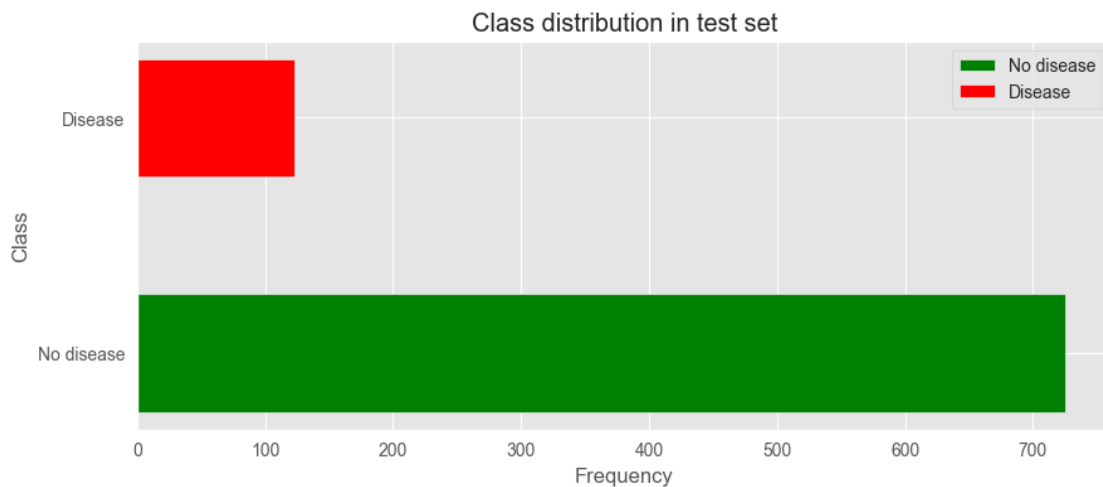
1.5.1 Class Distribution

```
[26]: values = y_train.value_counts()
indices = [k for k in values.index]
plt.figure(figsize=(10,4))
plt.title('Class distribution in train set')
plt.barh(indices, values, color=['green', 'red'], label=['No disease',
    ↪'Disease'], height=0.5)
plt.xlabel('Frequency')
plt.ylabel('Class')
plt.yticks(indices, ['No disease', 'Disease'])
```

```
plt.legend()
plt.show()
```



```
[27]: values = y_test.value_counts()
indices = [k for k in values.index]
plt.figure(figsize=(10,4))
plt.title('Class distribution in test set')
plt.barh(indices, values, color=['green', 'red'], label=['No disease', 'Disease'], height=0.5)
plt.xlabel('Frequency')
plt.ylabel('Class')
plt.yticks(indices, ['No disease', 'Disease'])
plt.legend()
plt.show()
```



1.5.2 Function to split the data into train and test sets

```
[28]: def split(df_features, df_target, test_size=0.2, random_state=42):  
        X_train, X_test, y_train, y_test = train_test_split(df_features, df_target,  
        ↪test_size=test_size, random_state=random_state)  
        return X_train, X_test, y_train, y_test
```

1.6 Resampling

SMOTE is an oversampling technique where the synthetic samples are generated for the minority class. This algorithm helps to overcome the overfitting problem posed by random oversampling.

As the dataset is imbalanced, we will use SMOTE to oversample the minority class.

```
[29]: X_train_smote, y_train_smote = SMOTE().fit_resample(X_train, y_train)
```

```
[30]: print("X_train_smote shape:", X_train_smote.shape)  
        print("y_train_smote shape:", y_train_smote.shape)
```

```
X_train_smote shape: (5742, 7)  
y_train_smote shape: (5742,)
```

```
[31]: values = y_train_smote.value_counts()  
        indices = [k for k in values.index]  
        plt.figure(figsize=(10, 4))  
        plt.title('Class distribution in training set after Resampling')  
        plt.barh(indices, values, color=['green', 'red'], label=['No disease',  
        ↪'Disease'], height=0.5)  
        plt.xlabel('Frequency')  
        plt.ylabel('Class')  
        plt.yticks(indices, ['No disease', 'Disease'])  
        plt.legend()  
        plt.show()
```



1.6.1 Function to perform resampling

```
[32]: def resample(X_train, y_train):  
        X_train_smote, y_train_smote = SMOTE().fit_resample(X_train, y_train)  
        return X_train_smote, y_train_smote
```

1.7 Model Pipelines

```
[33]: # Logistic Regression Model Pipeline  
logistic_pipeline = Pipeline([  
    ('Logistic Regression', LogisticRegression(solver='saga', max_iter=10000))  
])  
  
# SVM Model Pipeline  
linear_svc_pipeline = Pipeline([  
    ('Support Vector Machine', SVC())  
])  
  
# Decision Tree Model Pipeline  
decision_tree_pipeline = Pipeline([  
    ('Decision Tree', DecisionTreeClassifier())  
])  
  
# KNN Model Pipeline  
knn_pipeline = Pipeline([  
    ('K Nearest Neighbors', KNeighborsClassifier())  
])
```

```
pipelines = [logistic_pipeline, linear_svc_pipeline, decision_tree_pipeline,
             ↪knn_pipeline]
```

1.8 Modelling & Evaluation

```
[34]: # Train the models with cross validation
def train_model_cv(model_pipeline: Pipeline, X_train, y_train):
    cross_validate_results = cross_validate(model_pipeline, X_train, y_train,
    ↪cv=5, scoring=['accuracy', 'precision', 'recall', 'f1', 'roc_auc'],
    ↪return_train_score=True)
    return cross_validate_results
```

```
[35]: results = []
for pipeline in tqdm(pipelines):
    results.append(train_model_cv(pipeline, X_train_smote, y_train_smote))
```

100%| | 4/4 [00:08<00:00, 2.22s/it]

1.8.1 Results comparison of different models

```
[36]: table = []
for i in range(len(pipelines)):
    table.append([
        list(pipelines[i].named_steps.keys())[0],
        results[i]['test_accuracy'].mean(), results[i]['test_precision'].mean(),
        results[i]['test_recall'].mean(), results[i]['test_f1'].mean()
    ])
print(tabulate(table, headers=['Model', 'Accuracy', 'Precision', 'Recall',
    ↪'f1-score'], tablefmt='grid'))
```

Model	Accuracy	Precision	Recall	f1-score
Logistic Regression	0.664406	0.665482	0.660752	0.663071
Support Vector Machine	0.706549	0.68874	0.753751	0.719744
Decision Tree	0.762977	0.739813	0.812616	0.774393
K Nearest Neighbors	0.811564	0.736243	0.971438	0.837575

1.9 Inferencing

1.9.1 Problem Statement and Dataset Description

Heart disease is a major cause of death worldwide. The main objective of this project is to predict whether a person has heart disease or not based on the given features. The dataset contains 15 features and 1 target variable.

1.9.2 Preprocessing

The dataset is imbalanced. So, we have used SMOTE to oversample the minority class. PCA is used to reduce the number of features to 7. The dataset is then split into train and test sets.

1.9.3 Training

I have used 4 models: Logistic Regression, Support Vector Machine, Decision Tree and K-Nearest Neighbors. The models are trained using 5-fold cross validation with the following results of mean accuracy, precision, recall, f1 score and roc auc score:

Model	Accuracy	Precision	Recall	F1 Score	ROC AUC
Logistic Regression	0.673065	0.672859	0.673684	0.673227	0.728765
Support Vector Machine	0.70418	0.689422	0.743344	0.715261	0.771742
Decision Tree	0.763003	0.74574	0.799071	0.771323	0.763003
K Nearest Neighbors	0.811765	0.735244	0.974923	0.838154	0.917557

1.9.4 Strengths and weaknesses of the models

Logistic Regression:

- **Strengths:**
 - **Balanced Performance:** Logistic regression shows a relatively balanced performance across accuracy, precision, recall, and F1 score.
 - **Interpretability:** Coefficients of logistic regression provide insights into the importance of features.
- **Weaknesses:**
 - **Limited Predictive Power:** The model may not capture complex non-linear relationships effectively, reflected in slightly lower performance compared to other models.

Support Vector Machine (SVM):

- **Strengths:**
 - **High Recall:** SVM shows high recall, indicating its effectiveness in identifying true positives (patients with heart disease) accurately.
 - **Non-Linear Decision Boundary:** SVM effectively models non-linear relationships.
- **Weaknesses:**
 - **Complexity and Computation:** SVM tends to be computationally intensive, especially as the dataset size grows.
 - **Sensitivity to Hyperparameters:** Performance is sensitive to kernel choice and hyperparameter tuning.

Decision Tree:

- **Strengths:**
 - **High Recall:** Decision trees demonstrate a high recall, indicating effective identification of individuals with heart disease.
 - **Interpretability:** Easy to interpret and explain due to its hierarchical structure.
- **Weaknesses:**
 - **Overfitting:** Prone to overfitting, as reflected in the difference in performance between training and testing data.

K Nearest Neighbors (KNN):

- **Strengths:**
 - **High Recall:** KNN shows very high recall, suggesting it effectively identifies individuals with heart disease, minimizing false negatives.
 - **Non-Parametric and Lazy Learning:** KNN is non-parametric and adapts to any new data without retraining.
- **Weaknesses:**
 - **Computationally Intensive:** The model is computationally intensive, especially as the dataset size increases.
 - **Sensitivity to Distance Metric:** Performance can vary based on the choice of distance metric.

Overall Comparison:

- Decision Tree performs well, striking a balance between recall and precision. It is also interpretable, aiding understanding.
- Logistic Regression offers a good balance across metrics but may lack the predictive power compared to more complex models.
- SVM performs well, particularly in recall, but requires careful parameter tuning due to sensitivity.
- In this context, KNN stands out with very high recall, crucial for identifying individuals at risk of heart disease. However, it sacrifices precision, leading to a lower F1 score.

1.9.5 Best Model

The best model is K-Nearest Neighbors with an accuracy of 81.12% and roc auc score (efficiency) of 91.76%. For the purpose of detecting heart disease, we would want to minimize the false negatives, i.e., the number of people who have heart disease but are predicted to not have heart disease as it is more dangerous. So, we will choose the model with the highest recall, which is, again, K-Nearest Neighbors.

1.10 Apply Model

1.10.1 Loading and preprocessing the data

```
[37]: df_features, df_target = load_data()
      df_features = feature_selection(df_features)
      X_train, X_test, y_train, y_test = split(df_features, df_target)
      X_train, y_train = resample(X_train, y_train)
```

1.10.2 Training K-Nearest Neighbors model

```
[38]: model = KNeighborsClassifier()
```

```
[39]: history = model.fit(X_train, y_train)
```

1.10.3 Predicting and Evaluating the model on the test set

```
[40]: y_pred = model.predict(X_test)
```

```
[41]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.68	0.77	725
1	0.21	0.49	0.29	123
accuracy			0.65	848
macro avg	0.55	0.58	0.53	848
weighted avg	0.79	0.65	0.70	848

```
[42]: print(f"Accuracy: {accuracy_score(y_test, y_pred)*100:.2f}%")
      print(f"Precision: {precision_score(y_test, y_pred)*100:.2f}%")
      print(f"Recall: {recall_score(y_test, y_pred)*100:.2f}%")
      print(f"F1 Score: {f1_score(y_test, y_pred):.6f}")
      print(f"ROC AUC Score: {roc_auc_score(y_test, y_pred):.6f}")
```

Accuracy: 65.21%
Precision: 20.55%
Recall: 48.78%
F1 Score: 0.289157
ROC AUC Score: 0.583902

1.11 Saving the model

```
[43]: dump(model, open('model.pkl', 'wb'))
```